Reinforcement Learning With Hado van Hasselt

Yaro Kazakov

September 14, 2024

Abstract

These are study notes to prepare myself for PhD applications.

1 Lecture 3 - Finite Markov Decision Processes

We try to solve MDPs in the rest of this book. This problem involves evaluative feedback, as in bandits, but also an associative aspect—choosing different actions in different situations. MDPs are a classical formalization of sequential decision making, where actions influence not just immediate rewards, but also subsequent situations, or states, and through those future rewards. Thus MDPs involve **delayed reward** and the need to **tradeoff immediate and delayed reward**. Whereas in bandit problems we estimated the value $q_*(a)$ of each action a, in MDPs we estimate the value $q_*(s, a)$ of each action a in each state s, or we estimate the value $v_*(s)$ of each state given optimal action selections.

As in all of artificial intelligence, there is a tension between **breadth of applicability and mathematical tractability**.

1.1 The Agent-Environment Interface

MDPs are meant to be a straightforward framing of the problem of learning from interaction to achieve a goal. The learner and decision maker is called the agent. The thing it interacts with, comprising everything outside the agent, is called the environment.
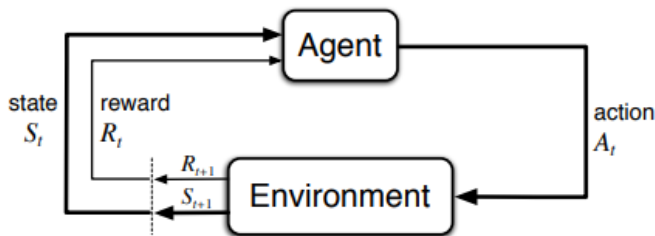

Figure 3.1: The agent–environment interaction in a Markov decision process.

Figure 1: Agent-environment Interaction

The MDP and agent together thereby give rise to a sequence or trajectory that begins like this:
$S_0, A_0, R_1, S_1, A_1, R_2, S_2, A_2, ...$

In a finite MDP, the sets of states, actions, and rewards (S, A, and R) all have a FINITE number of elements.

That is, there is a probability of those values occurring at time t, given particular values of the preceding state and action:

$p(s', r|s, a) \doteq Pr S_t = s', R_t = r|S_{t-1} = s, A_{t-1} = a$, the function p defines the **dynamics of the MDP**. The dot over the equals sign means it is a definition. Here this reminds us that:

$\sum_{s' \in S} \sum_{r \in R} p(s', r|s, a) = 1$, for all $s \in S$, $a \in A(s)$

In MDP, the probabilities given by p completely characterize the environment's dynamics. The state must include information about all aspects of the past agent–environment interaction that make a difference for the future. If it does, then the state is said to have the Markov property.

Markov property is assumed throughout this book.

For the **state transition probabilities** (not joint distribution), we do the following:

$p(s'|s, a) \doteq Pr S_t = s'|S_{t-1} = s, A_{t-1} = a = \sum_{r \in R} p(s', r|s, a)$, why is that?

This is a concept of **marginalization**. Marginalization in probability theory refers to **the process of summing (or integrating) over the possible values of one or more variables** to obtain the probability distribution of a subset of the variables. The goal is to "marginalize" out (i.e., eliminate) certain variables from a joint distribution to focus on the probability distribution of the remaining variables.

We can also compute state-action-next-state: $r(s, a) \doteq E[Rt|S_{t-1} = s, A_{t-1} = a] = \sum_{r \in R} r \sum_{s \in S} (s', r|s, a)$

The general rule we follow is that anything that cannot be changed arbitrarily by the agent is considered to be outside of it and thus part of its environment.

1.1.1 Exercise 3.1

Exercise 3.1 Devise three example tasks of your own that fit into the MDP framework, identifying for each its states, actions, and rewards. Make the three examples as different from each other as possible. The framework is abstract and flexible and can be applied in many different ways. Stretch its limits in some way in at least one of your examples.

1. Walk from a bedroom to a kitchen. States: in a bedroom, walking to the kitchen, in the kitchen. Actions: Walking, standing still. Rewards: 0, -1.

2. Room air-conditioning. State: Hot, OK, cold. Action: Cool down, warm up, maintain, idle. Reward: -1,-1,1,0.

3. Cruise control (assuming no cars and no traffic). States: Too fast, too slow, OK. Action: Accelerate, decelerate, maintain. Rewards: -1, -1, 0.

1.1.2 Exercise 3.2

Is the MDP framework adequate to usefully represent all goal-directed learning tasks? Can you think of any clear exceptions?

1. No, a situation where states are not self-containing cannot be modelled with MDPs. For example a robot that is controled through a virtual camera cannot send action signals to the robot based on one snapshot. The snapshot might relate to a situation where a robot is charging forward, hopping or turning. At some moment time t, this snapshot would be insufficient to judge whether a robot is doing any of the three actions. Hence, state history would be required.

1.1.3 Exercise 3.3

Consider the problem of driving. You could define the actions in terms of the accelerator, steering wheel, and brake, that is, where your body meets the machine. Or you could define them farther out—say, where the rubber meets the road, considering your actions to be tire torques. Or you could define them farther in—say, where your brain meets your body, the actions being muscle twitches to control your limbs. Or you could go to a really high level and say that your actions are your choices of where to drive. What is the right level, the right place to draw the line between agent and environment? On what basis is one location of the line to be preferred over another? Is there any fundamental reason for preferring one location over another, or is it a free choice?

1. High-level abstraction would be navigation. Actions are route choices, states are your position on the map, reward is the distance differential between your position and destination or time to get there. Every additional minute is a penalty of -1.

2. Autonomous driving algo is a low-level abstraction. Actions are steering, accelerating and braking.

1.1.4 Exercise 3.4

The joint probability would be the same because for every s,a,s' triple there is onle one deterministic reward. Use the formula 3.6

1.2 Returns and Episodes

We seek to maximize expected return. Return is defined as before:

$G_t = R_{t+1} + R_{t+2} + \cdots + R_T$

This notion works if there is a **final step**. That is when the agent-environment interaction breaks naturally into subsequences, called **episodes**. Each episode ends in a special terminal state.
On the other hand, an ongoing task of a a robot with a life long span would be called **continuous tasks**.

Discounting is also an important concept. $G_t \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+2} + \ldots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$ where $\gamma is between 0 and 1$

$G_t = R_{t+1} + \gamma G_{t+1}$, **this recursive formula is very important for RL**.

Note that even though the return is infinite, it is still finite if the reward is nonzero and constant. if $\gamma < 1$. If $R_t = +1$

$G_t = 1/(1 - \gamma)$

1.2.1 Exercise 3.5

S becomes S^+ because we need to include the terminal space explicitly.

1.2.2 Exercise 3.8

$G_5 = 0, G_4 = 2 + 0.5 * G_5 = 2, G_3 = 3 + 0.25 * G4 = 3.5$ etc..

1.2.3 Exercise 3.9

Same as before but using sum to infinity.

1.3 Policies and Value Functions

State value function under/for policy π $v_{pi}(s) \doteq E_\pi[G_t|S_t = s] = E_\pi[\sum_{k=0}^{\inf} \gamma^k R_{t+k+1}|S_t = s]$ Action-value function for policy π: $q_\pi(s, a) \doteq E_\pi[G_t|S_t, A_t = a] = E_\pi[\sum_{k=0}^{\inf} \gamma^k R_{t+k+1}|S_t = s, A_t = a]$

1.3.1 Exercise 3.12

Exercise 3.12 Give an equation for v in terms of q and pi.

This is yet another marginalization Rule $P(X|Y) = \sum_{z \in Z} P(X|Y,Z) * P(Z|Y)$ the state value function is the weighted sum of the state action value function, with the weights equal to the probabilities of selecting each action.

$$v_\pi(s) = \sum_a \pi(a|s) * q(s,a)$$

1.3.2 Exercise 3.13

Just like Bellman but without the first term, i.e. averaging by the probability of taking action a given s.

1.4 Bellman equation for v_π

$v_\pi(s) = \sum_a \pi(a,s) \sum_{(s',r)} p(s',r|s,a)[r + \gamma v_\pi(s')]$ for all s ∈ S.

The Bellman equation (3.14) averages over all the possibilities, weighting each by its probability of occurring. It states that the value of the start state must equal the (discounted) value of the expected next state, plus the reward expected along the way.

1.4.1 Exercise 3.14

P(a—s) is equal to 1/4, four-term probability is 1, because the state is fully deterministic once the action is taken. r is equal to 0. Plug in the numbers from the near cells and multiply by 0.9.

1.4.2 Exercise 3.15

The intervals between them. Actions are chosen on a relative basis, if an action-value function is higher relative to all the other functions then it will be chosen. $G_t + c = R_{t+1} + c + \gamma(R_{t+2} + c) + \gamma^2(R_{t+3} + c) + \ldots$

$$G'_t = G_t + (c + \gamma * c\gamma^2 * c + \ldots)$$

$$G'_t = G_t + c/(1 - \gamma)$$

meaning that the total reward shifts by $c/(1 - \gamma) = v_c$(like in the question)

1.4.3 Exercise 3.16

Now consider adding a constant c to all the rewards in an episodic task, such as maze running. Would this have any effect, or would it leave the task unchanged as in the continuing task above? Why or why not? Give an example.

Now the effect is different, since the sum of rewards is finite. $G_t + c = R_{t+1} + c + \gamma(R_{t+2} + c) + \gamma^2(R_{t+3} + c) + \cdots + \gamma^{T-t-1}R_T$

$$G'_t = G_t + c(1 - \gamma^(T - t))/(1 - \gamma)$$

So, the shift affects different cumulative rewards differently depending on $T - t$

Assume a game with two actions and one State. $A_0 = 0, A_1 = 1$. A_0 takes the agent to state S, A_1 terminates and returns 1. If we shift the rewards by 1 and apply a discount factor of 99/100 (minimal discount) we get that instead of terminating early and receiving the reward +1 immediately, the agent will be stuck at action A_0 to get an additional $1/(1 - \gamma)$ in infinity.

1.4.4 Exercise 3.17

Fun question. Using 3.14 we can deduce that

$$q_\pi(s) = \sum_{r \in R, s' \in S} p(s', r|s, a)[r + \gamma v_\pi(s) = \sum_{r \in R, s' \in S} p(s', r|s, a)[r + \gamma \sum_{a'} \pi(a'|s)(q(a', s'))$$

1.4.5 Exercise 3.18

Average the action functions over the policies $v_\pi(s) = E_\pi[G_t|S_t = s] = E_\pi[q_\pi(s, a)] \sum_a \pi(a|s) * q_\pi(a, s)$

1.4.6 Exercise 3.19

$$q_\pi(a, s) = E_\pi[R_{t+1} + \gamma G_{t+1}|S_t = s, A_t = a] = \sum_{r \in R, s' \in S} p(s', r|s, a)[r + \gamma v_\pi(s')]$$

1.5 Bellman Optimal Equation

$$v_*(s) = max_{a \in A(s)} q_{\pi_*}(s, a) = max_a \sum_{s', r} p(s', r|s, a)[r + \gamma v_*(s')]$$

For each states, there will be one or more actions at which the maximum is obtained in the Bellman optimality equation. Any policy that assigns nonzero probability only to these actions is an optimal policy.

Solving a reinforcement learning task means, roughly, finding a policy that achieves a lot of reward over the long run. For finite MDPs, we can precisely define an optimal policy in the following way. Value functions define a partial ordering over policies. A policy π is defined to be better than or equal to a policy π' if its expected return is greater than or equal to that of π' for all states.

$$q_*(s, a) = \sum_{r \in R, s' \in S} p(s', r|s, a)[r + \gamma v_*(s')] = \sum_{r \in R, s' \in S} p(s', r|s, a)[r + \gamma max_{a'} q_*(s', a')]$$

Optimal policies also share the same optimal action-value function, denoted q_*, and defined as

$$q_*(s, a) \doteq max_\pi q_\pi(s, a)$$

1.5.1 Exercise 3.20

Off the green - use driver, on the green use putter. The optimal state value from the tee is -3.

1.5.2 Exercise 3.21

Same as before? The diagram is the same as the diagram for v_{putt}

1.5.3 Exercise 3.22

Left, both, right.

1.5.4 Exercise 3.23

$$q_*(s, a) = \sum_{r \in R, s' \in S} p(s', r|s, a)[r + \gamma v_*(s')]$$

$$q_*(s, a) = \sum_{r \in R, s' \in S} p(s', r|s, a)[r + \gamma argmax_{a'} q_*(s', a')]$$

1.5.5 Exercise 3.24

Figure 3.5 gives the optimal value of the best state of the gridworld as 24.4, to one decimal place. Use your knowledge of the optimal policy and (3.8) to express this value symbolically, and then to compute it to three decimal places

$$v_*(s) = \sum_{r \in R, s' \in S} p(s', r|s, a)[r + \gamma v_\pi(s')]$$

$$v_*(s) = 10 + 0.9 * 16$$

1.5.6 Exercise 3.25

Give an equation for v* in terms of q*.

$$v_*(s) = \sum_a \pi_*(a|s) * (q_*(a, s)$$

1.5.7 Exercise 3.26

Give an equation for q* in terms of v* and the four-argument p.

Equation 3.17 but with the four-parameter p instead of the expectations.

1.5.8 Exercise 3.27

Give an equation for π_* in terms of q*.

Kind of like a greedy algorithm? We sum over the actions that are greedy and divide by the number of all actions at a given state?

$$\pi_*(a|s) = (greedy a)/(number of actions at a state)$$

1.5.9 Exercise 3.28

Give an equation for $\pi_*(a|s)$ in terms of v_* and the four-argument p.

Take the answer in 3.27 and plug in 3.26